

Max-flow and min-cut problems

Edmonds
Karp alg
Generic
reduction to
MaxFlow



1 Edmonds Karp alg

2 Generic reduction to MaxFlow

Dinic and Edmonds-Karp algorithm

Edmonds
Karp alg

Generic
reduction to
MaxFlow

J. Edmonds, R. Karp: *Theoretical improvements in algorithmic efficiency for network flow problems*. Journal ACM 1972.

Yefim Dinic: *Algorithm for solution of a problem of maximum flow in a network with power estimation*. Doklady Ak.N. 1970

Choosing a **good**
augmenting path can lead
to a faster algorithm.

Use **BFS** to find an
augmenting paths in G_f .



Edmonds-Karp algorithm

Edmonds
Karp alg

Generic
reduction to
MaxFlow

FF algorithm but **using BFS**: choose the augmenting path in G_f with the smallest length (number of edges).

Edmonds-Karp(G, c, s, t)

For all $e = (u, v) \in E$ let $f(u, v) = 0$

$G_f = G$


while there is an $s \rightsquigarrow t$ path in G_f
do

$P = \text{BFS}(G_f, s, t)$

$f = \text{Augment}(f, P)$

 Compute G_f

return f



The BFS in EK will
choose:  or 

BFS paths on G_f

Edmonds
Karp alg

Generic
reduction to
MaxFlow

For $\mathcal{N} = (V, E, c, s, t)$ and a flow f in \mathcal{N} , assuming that G_f has an augmenting path, let f' be the next flow after executing one step of the EK algorithm.

- The path from s to t in a BFS traversal starting at s , is a path $s \rightsquigarrow t$ with minimum number of edges, i.e., a **shortest length path**.
- For $v \in V$, let $\delta_f(s, v)$ denote length of a shortest length path from s to v in G_f .

Some properties of G_f and $G_{f'}$

Edmonds
Karp alg

Generic
reduction to
MaxFlow

How can we have $(u, v) \in E_{f'}$ but $(u, v) \notin E_f$?

- (u, v) is a forward edge saturated in f and not in f' .
- (u, v) is a backward edge in G_f and $f(v, u) = 0$

In any of the two cases, the augmentation must have modified the flow from v to u , so (u, v) must form part of the augmenting path.

EK and the shortest length distances

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Lemma

If the EK-algorithm runs on $\mathcal{N} = (V, E, c, s, t)$, for all vertices $v \neq s$, $\delta_f(s, v)$ increases monotonically with each flow augmentation.

Proof. By contradiction.

Let f be the first flow such that, for some $u \neq s$,

$$\delta_{f'}(s, u) < \delta_f(s, u).$$

EK and the shortest length distances

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Proof (cont)

Let v be the vertex with the minimum $\delta_{f'}(s, v)$ whose distance was decreased.

- Let $P : s \rightsquigarrow u \rightarrow v$ be a shortest length path from s to v in $G_{f'}$
- Then, $\delta_{f'}(s, v) = \delta_{f'}(s, u) + 1$ and $\delta_{f'}(s, u) \geq \delta_f(s, u)$.
- If $(u, v) \in E_f$,
 $\delta_f(s, v) \leq \delta_f(s, u) + 1 \leq \delta_{f'}(s, u) + 1 = \delta_{f'}(s, v)$
- So, $(u, v) \notin E_f$

EK and the shortest length distances

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Proof (cont)

How can we have ?

- $(u, v) \in E_{f'}$ but $(u, v) \notin E_f$
- If so, (v, u) appears in the augmenting path.
- Then, the shortest length path from s to u in G_f has (v, u) as its last edge.
$$\delta_f(s, v) \leq \delta_f(s, u) - 1 \leq \delta_{f'}(s, u) - 1 = \delta_{f'}(s, v) - 1 - 1$$
- which contradicts $\delta_{f'}(s, v) < \delta_f(u, v)$. □

Some properties of G_f and $G_{f'}$

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Let P be an augmenting path in G_f .

$(u, v) \in P$ is **critical** if $b(P) = c_f(u, v)$.

Critical edges do not appear in $G_{f'}$.

- (u, v) forward, $f'(u, v) = c(u, v)$
- (u, v) backward, $f'(v, u) = 0$

EK and critical edges

Lemma

In the EK algorithm, each one of the edges can become critical at most $|V|/2$ times.

Proof:

- Let $(u, v) \in E$, when (u, v) is critical for the first time,
 $\delta_f(s, v) = \delta_f(s, u) + 1$
- After this step (u, v) disappears from the residual graph until after the flow in (u, v) changes.
- At this point, (v, u) forms part of the augmenting path in $G_{f'}$, and $\delta_{f'}(s, u) = \delta_{f'}(s, v) + 1$,

$$\delta_{f'}(s, u) = \delta_{f'}(s, v) + 1 \geq \delta_f(s, v) + 1 \geq \delta_f(s, u) + 2$$

- So, the distance has increased by at least 2.



Complexity of Edmonds-Karp algorithm

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Theorem

The EK algorithm runs in $O(mn(n + m))$ steps. Therefore it is a polynomial time algorithm.

Proof:



- Need time $O(m + n)$ to find the augmenting path using BFS.
 - By the previous Lemma, there are $O(mn)$ augmentations.
-

Finding a min-cut

Given (G, s, t, c) to find a min-cut:

- 1 Compute the max-flow f^* in G .
- 2 Obtain G_{f^*} .
- 3 Find the set $S = \{v \in V \mid s \rightsquigarrow v\}$ in G_{f^*} .
- 4 Output the cut
 $(S, V - \{S\}) = \{(v, u) \mid v \in S \text{ and } u \in V - \{S\}\}$ in G .

The running time is the same than the algorithm to find the max-flow.



The max-flow problems: History

Edmonds
Karp alg

Generic
reduction to
MaxFlow

- Ford-Fulkerson (1956) $O(mC)$, where C is the max flow.
- Dinic (1970) (blocking flow) $O(n^2m)$
- Edmond-Karp (1972) (shortest augmenting path) $O(nm^2)$
- Karzanov (1974), $O(n^2m)$ Goldberg-Tarjant (1986) (push re-label preflow + dynamic trees) $O(nm \lg(n^2/m))$ (uses parallel implementation)
- King-Rao-Tarjan (1998) $O(nm \log_{m/n} n)$.
- J. Orlin (2013) $O(nm)$ (clever follow up to KRT-98)

So: Maximum flows can be computed in $O(nm)$ time!

Edmonds
Karp alg

Generic
reduction to
MaxFlow

1 Edmonds Karp alg

2 Generic reduction to MaxFlow

Applications: Generalized assignment problems

- Consider a **generalized assignment problem** \mathcal{GP} where, we have as input d finite sets X_1, \dots, X_d , each representing a different set of resources.
- Our goal is to choose the "largest" number of d -tuples, each d -tuple containing exactly one element from each X_i , subject to the constraints:
 - For each $i \in [d]$, each $x \in X_i$ can appear in at most $c(x)$ selected tuples.
 - For each $i \in [d]$, any two $x \in X_i$ and $y \in X_{i+1}$ can appear in at most $c(x, y)$ selected tuples.
 - The values for $c(x)$ and $c(x, y)$ are either in \mathbb{Z}^+ or ∞ .
- Notice that only pairs of objects between adjacent X_i and X_{i+1} are constrained.

Applications: Generic reduction to Max-Flow

Edmonds
Karp alg


Generic
reduction to
MaxFlow

Make the reduction from \mathcal{GP} to the following network \mathcal{N} :

- V contains a vertex x , for each element x in each X_i , and a copy x' , for each element $x \in X_i$ for $1 \leq i < d$.
- We add vertex s and vertex t .
- Add an edge $s \rightarrow x$ for each $x \in X_1$ and add an edge $y \rightarrow t$ for every $y \in X_d$. Give capacities $c(s, x) = c(x)$ and $c(y, t) = c(y)$.
- Add an edge $x' \rightarrow y$ for every pair $x \in X_i$ and $y \in X_{i+1}$. Give a capacity $c(x, y)$. Omit the edges with capacity 0.
- For every $x \in X_i$ for $1 \leq i < d$, add an edge $x \rightarrow x'$ with $c(x, x') = c(x)$.

Every path $s \rightsquigarrow t$ in \mathcal{N} identifies a feasible d -tuple, conversely every d -tuple determines a path $s \rightsquigarrow t$.

Flow Network: The reduction



- To solve \mathcal{GP} , we construct \mathcal{N} , and then we find an integer maximum flow f^* .
- In the subgraph formed by edges with $f^*(e) > 0$, we find a (s, t) path P (a d -tuple), decrease in 1 the flow in each edge of P , remove edges with 0 flow.
- We repeat the procedure for $|f^*|$ times. In this way we obtain a set of d -tuples with maximum size verifying all the restrictions.

Edmonds
Karp alg

Generic
reduction to
MaxFlow

FINAL'S SCHEDULING

We have as input:

- n courses, each one with a final. Each exam must be given in one room. Each course c_i has $E[i]$ students.
- r rooms. Each r_j has a capacity $S[j]$,
- τ time slots. For each room and time slot, we only can schedule one final.
- p professors to watch exams. Each exam needs one professor in each class and time. Each professor has its own restrictions of availability and no professor can oversee more than 6 finals. For each p_ℓ and τ_k define a Boolean variable $A[k, \ell] = T$ if p_ℓ is available at τ_k .

Design an efficient algorithm that correctly schedules a room, a time slot and a professor to every final, or report that not such schedule is possible.

Construction of the network

Edmonds
Karp alg

Generic
reduction to
MaxFlow

Construct the network \mathcal{N} with vertices


$\{s, t, \{c_i\}, \{r_j\}, \{\tau_k\}, \{p_\ell\}\}$. Edges and capacities:

- (s, c_i) with capacity 1 (each course has one final)
- (c_i, r_j) , if $E[i] \leq S[j]$, with capacity ∞
- $\forall j, k, (r_j, \tau_k)$, with capacity 1 (one final per room and time slot).
- (τ_k, p_ℓ) , if $A[k, \ell] = T$, capacity 1
(p can watch one final, if p is available at τ_k).
- (p_ℓ, t) , capacity 6 (each p can watch ≤ 6 finals)

Notice that neither rooms nor time slots have individual restrictions.

FINAL'S SCHEDULING: Flow Network

Edmonds
Karp alg
Generic
reduction to
MaxFlow



FINAL'S SCHEDULING

Edmonds
Karp alg

Generic
reduction to
MaxFlow

- Notice the input size to the problem is $N = n + r + \tau + p + 2$. and size of the network is $O(N)$ vertices and $O(N^2)$ edges, why?
- Every path $s \rightsquigarrow t$ is an assignment of room-time-professor to a final, and any assignment room-time-professor to a final can be represented by a path $s \rightsquigarrow t$.
- Every integral flow identifies a collection of $|f|$ (s, t) -paths leading to a valid assignment for $|f|$ finals and viceversa.

FINAL'S SCHEDULING

Edmonds
Karp alg
Generic
reduction to
MaxFlow

- To maximize the number of finals to be given, we compute the max-flow f^* from s to t .
- If $|f^*| = n$, then we can schedule all finals, otherwise we can not.
- To recover the assignment we have to consider the edges with positive flow and extract assignment from the n (s, t) -paths
- **Complexity:**
 - To construct \mathcal{N} , we need $O(N^2)$.
 - As $|f^*| \leq n$ integral, we can use Ford-Fulkerson to compute f^* , with cost $O(nN^2)$.
 - The second part requires $O(N^2)$ time.
 - So, the cost of the algorithm is $O(nN^2) = O(n(n + r + \tau + p)^2)$.